

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
# Necessary Imports
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("/content/drive/MyDrive/Final-Year AI/week2/mnist_dataset.csv") # changed to read_csv and the correct file name
# Step 2: Dataset Information
print("Dataset Preview:")
print(df.head()) # Show first 5 rows
print("\nDataset Information:")
print(df.info()) # Summary of dataset
```

Dataset Preview:

	label	pixel_0	pixel_1	pixel_2	pixel_3	pixel_4	pixel_5	pixel_6	\
0	5	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	4	0	0	0	0	0	0	0	
3	1	0	0	0	0	0	0	0	
4	9	0	0	0	0	0	0	0	

	pixel_7	pixel_8	...	pixel_774	pixel_775	pixel_776	pixel_777	\
0	0	0	...	0	0	0	0	
1	0	0	...	0	0	0	0	
2	0	0	...	0	0	0	0	
3	0	0	...	0	0	0	0	
4	0	0	...	0	0	0	0	

	pixel_778	pixel_779	pixel_780	pixel_781	pixel_782	pixel_783
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

[5 rows x 785 columns]

```
Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60000 entries, 0 to 59999
Columns: 785 entries, label to pixel_783
dtypes: int64(785)
memory usage: 359.3 MB
None
```

## ✓ Some Helper Function:

### ✓ Softmax Function:

```
import numpy as np
```

```
def softmax(z):
    """
    Compute the softmax probabilities for a given input matrix.

    Parameters:
    z (numpy.ndarray): Logits (raw scores) of shape (m, n), where
        - m is the number of samples.
        - n is the number of classes.

    Returns:
    numpy.ndarray: Softmax probability matrix of shape (m, n), where
        each row sums to 1 and represents the probability
        distribution over classes.

    Notes:
    - The input to softmax is typically computed as: z = XW + b.
```

```

- Uses numerical stabilization by subtracting the max value per row.
"""

# Your Code Here.
# Subtract max for numerical stability
z_stable = z - np.max(z, axis=1, keepdims=True)

# Compute exponentials
exp_z = np.exp(z_stable)

# Normalize by sum of exponentials per row
return exp_z / np.sum(exp_z, axis=1, keepdims=True)

return

```

## ✓ Softmax Test Case:

This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax.

```

# Example test case
z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])
softmax_output = softmax(z_test)

# Verify if the sum of probabilities for each row is 1 using assert
row_sums = np.sum(softmax_output, axis=1)

# Assert that the sum of each row is 1
assert np.allclose(row_sums, 1), f"Test failed: Row sums are {row_sums}"

print("Softmax function passed the test case!")

```

→ Softmax function passed the test case!

## ✓ Prediction Function:

```

def predict_softmax(X, W, b):
    """
    Predict the class labels for a set of samples using the trained softmax model.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.
    W (numpy.ndarray): Weight matrix of shape (d, c), where c is the number of classes.
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    numpy.ndarray: Predicted class labels of shape (n,), where each value is the index of the predicted class.
    """
    # predicted_classes = # Your Code Here

    # Compute logits (raw scores)
    logits = np.dot(X, W) + b # Shape: (n, c)

    # Compute softmax probabilities
    probabilities = softmax(logits)

    # Select the class with the highest probability
    predicted_classes = np.argmax(probabilities, axis=1)

    return predicted_classes

```

## ✓ Test Function for Prediction Function:

The test function ensures that the predicted class labels have the same number of elements as the input samples, verifying that the model produces a valid output shape.

```

# Define test case
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

```

```
# Expected Output:
# The function should return an array with class labels (0, 1, or 2)

y_pred_test = predict_softmax(X_test, W_test, b_test)

# Validate output shape
assert y_pred_test.shape == (3,), f"Test failed: Expected shape (3,), got {y_pred_test.shape}"

# Print the predicted labels
print("Predicted class labels:", y_pred_test)
```

↗ Predicted class labels: [1 1 0]

## ▼ Loss Function:

```
import numpy as np

def loss_softmax(y_pred, y):
    """
    Compute the cross-entropy loss for a single sample.

    Parameters:
    y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single sample,
                           where c is the number of classes.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is the number of classes.

    Returns:
    float: Cross-entropy loss for the given sample.
    """
    epsilon = 1e-12 # Small value to prevent log(0)
    y_pred = np.clip(y_pred, epsilon, 1.0 - epsilon) # Clipping for numerical stability
    loss = -np.sum(y * np.log(y_pred))

    return loss
```

## ▼ Test case for Loss Function:

This test case Compares loss for correct vs. incorrect predictions.

- Expects low loss for correct predictions.
- Expects high loss for incorrect predictions.

```
import numpy as np

# Define correct predictions (low loss scenario)
y_true_correct = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True one-hot labels
y_pred_correct = np.array([[0.9, 0.05, 0.05],
                           [0.1, 0.85, 0.05],
                           [0.05, 0.1, 0.85]]) # High confidence in the correct class

# Define incorrect predictions (high loss scenario)
y_pred_incorrect = np.array([[0.05, 0.05, 0.9], # Highly confident in the wrong class
                             [0.1, 0.05, 0.85],
                             [0.85, 0.1, 0.05]])

# Compute loss for both cases
loss_correct = loss_softmax(y_pred_correct, y_true_correct)
loss_incorrect = loss_softmax(y_pred_incorrect, y_true_correct)

# Validate that incorrect predictions lead to a higher loss
assert loss_correct < loss_incorrect, f"Test failed: Expected loss_correct < loss_incorrect, but got {loss_correct:.4f} >= {loss_incorrect:.4f}"

# Print results
print(f"Cross-Entropy Loss (Correct Predictions): {loss_correct:.4f}")
print(f"Cross-Entropy Loss (Incorrect Predictions): {loss_incorrect:.4f}")
```

↗ Cross-Entropy Loss (Correct Predictions): 0.4304  
Cross-Entropy Loss (Incorrect Predictions): 8.9872

## ✓ Cost Function:

```
import numpy as np
```

```
def cost_softmax(X, y, W, b):  
    """  
    Compute the average softmax regression cost (cross-entropy loss) over all samples.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.  
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where n is the number of samples and c is the number of classes.  
    W (numpy.ndarray): Weight matrix of shape (d, c).  
    b (numpy.ndarray): Bias vector of shape (c,).  
  
    Returns:  
    float: Average softmax cost (cross-entropy loss) over all samples.  
    """  
    n = X.shape[0] # Number of samples  
  
    # Compute logits (XW + b)  
    logits = np.dot(X, W) + b  
  
    # Apply softmax  
    exp_logits = np.exp(logits - np.max(logits, axis=1, keepdims=True)) # Numerical stability  
    y_pred = exp_logits / np.sum(exp_logits, axis=1, keepdims=True)  
  
    # Compute cross-entropy loss  
    epsilon = 1e-12 # Prevent log(0)  
    y_pred = np.clip(y_pred, epsilon, 1 - epsilon) # Clipping for numerical stability  
    total_loss = -np.sum(y * np.log(y_pred))  
  
    # Return average loss  
    return total_loss / n
```

## ✓ Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

```
import numpy as np
```

```
# Example 1: Correct Prediction (Closer predictions)  
X_correct = np.array([[1.0, 0.0], [0.0, 1.0]]) # Feature matrix for correct predictions  
y_correct = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, matching predictions)  
W_correct = np.array([[5.0, -2.0], [-3.0, 5.0]]) # Weights for correct prediction  
b_correct = np.array([0.1, 0.1]) # Bias for correct prediction  
  
# Example 2: Incorrect Prediction (Far off predictions)  
X_incorrect = np.array([[0.1, 0.9], [0.8, 0.2]]) # Feature matrix for incorrect predictions  
y_incorrect = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, incorrect predictions)  
W_incorrect = np.array([[0.1, 2.0], [1.5, 0.3]]) # Weights for incorrect prediction  
b_incorrect = np.array([0.5, 0.6]) # Bias for incorrect prediction  
  
# Compute cost for correct predictions  
cost_correct = cost_softmax(X_correct, y_correct, W_correct, b_correct)  
  
# Compute cost for incorrect predictions  
cost_incorrect = cost_softmax(X_incorrect, y_incorrect, W_incorrect, b_incorrect)  
  
# Check if the cost for incorrect predictions is greater than for correct predictions  
assert cost_incorrect > cost_correct, f"Test failed: Incorrect cost {cost_incorrect} is not greater than correct cost {cost_correct}"  
  
# Print the costs for verification  
print("Cost for correct prediction:", cost_correct)  
print("Cost for incorrect prediction:", cost_incorrect)  
  
print("Test passed!")
```

```
➡ Cost for correct prediction: 0.0006234364133349324  
Cost for incorrect prediction: 0.29930861359446115  
Test passed!
```

## ✓ Computing Gradients:

```
import numpy as np

def compute_gradient_softmax(X, y, W, b):
    """
    Compute the gradients of the cost function with respect to weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    tuple: Gradients with respect to weights (d, c) and biases (c,).
    """
    n = X.shape[0] # Number of samples

    # Compute logits (XW + b)
    logits = np.dot(X, W) + b

    # Apply softmax
    exp_logits = np.exp(logits - np.max(logits, axis=1, keepdims=True)) # Numerical stability
    y_pred = exp_logits / np.sum(exp_logits, axis=1, keepdims=True)

    # Compute gradient of loss w.r.t logits
    dL_dlogits = y_pred - y # Gradient of cross-entropy loss w.r.t logits

    # Compute gradients
    grad_W = np.dot(X.T, dL_dlogits) / n # Gradient w.r.t. weights
    grad_b = np.sum(dL_dlogits, axis=0) / n # Gradient w.r.t. biases

    return grad_W, grad_b
```

## ✓ Test case for compute\_gradient function:

The test checks if the gradients from the function are close enough to the manually computed gradients using `np.allclose`, which accounts for potential floating-point discrepancies.

```
import numpy as np

# Define a simple feature matrix and true labels
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
y_test = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True labels (one-hot encoded, 3 classes)

# Define weight matrix and bias vector
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

# Compute the gradients using the function
grad_W, grad_b = compute_gradient_softmax(X_test, y_test, W_test, b_test)

# Manually compute the predicted probabilities (using softmax function)
z_test = np.dot(X_test, W_test) + b_test
y_pred_test = softmax(z_test)

# Compute the manually computed gradients
grad_W_manual = np.dot(X_test.T, (y_pred_test - y_test)) / X_test.shape[0]
grad_b_manual = np.sum(y_pred_test - y_test, axis=0) / X_test.shape[0]

# Assert that the gradients computed by the function match the manually computed gradients
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W are not equal.\nExpected: {grad_W_manual}\nGot: {grad_W}"
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b are not equal.\nExpected: {grad_b_manual}\nGot: {grad_b}"

# Print the gradients for verification
print("Gradient w.r.t. W:", grad_W)
print("Gradient w.r.t. b:", grad_b)
```

```
print("Test passed!")
```

```
↩ Gradient w.r.t. W: [[ 0.1031051  0.01805685 -0.12116196]
 [-0.13600547  0.00679023  0.12921524]]
Gradient w.r.t. b: [-0.03290036  0.02484708  0.00805328]
Test passed!
```

## ✓ Implementing Gradient Descent:

```
import numpy as np

def gradient_descent_softmax(X, y, W, b, alpha, n_iter, show_cost=False):
    """
    Perform gradient descent to optimize the weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).
    alpha (float): Learning rate.
    n_iter (int): Number of iterations.
    show_cost (bool): Whether to display the cost at intervals.

    Returns:
    tuple: Optimized weights, biases, and cost history.
    """
    cost_history = []

    for i in range(n_iter):
        # Compute gradients
        grad_W, grad_b = compute_gradient_softmax(X, y, W, b)

        # Update weights and biases using gradient descent
        W -= alpha * grad_W
        b -= alpha * grad_b

        # Compute cost and store it
        cost = cost_softmax(X, y, W, b)
        cost_history.append(cost)

        # Display cost if required
        if show_cost and i % 100 == 0:
            print(f"Iteration {i}: Cost = {cost:.6f}")

    return W, b, cost_history
```

## ✓ Preparing Dataset:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def load_and_prepare_mnist(csv_file, test_size=0.2, random_state=42):
    """
    Reads the MNIST CSV file, splits data into train/test sets, and plots one image per class.

    Arguments:
    csv_file (str) : Path to the CSV file containing MNIST data.
    test_size (float) : Proportion of the data to use as the test set (default: 0.2).
    random_state (int) : Random seed for reproducibility (default: 42).

    Returns:
    X_train, X_test, y_train, y_test : Split dataset.
    """

    # Load dataset
    df = pd.read_csv(csv_file)

    # Separate labels and features
    y = df.iloc[:, 0].values # First column is the label
```

```

X = df.iloc[:, 1:].values # Remaining columns are pixel values

# Normalize pixel values (optional but recommended)
X = X / 255.0 # Scale values between 0 and 1

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=random_state)

# Plot one sample image per class
plot_sample_images(X, y)

return X_train, X_test, y_train, y_test

def plot_sample_images(X, y):
    """
    Plots one sample image for each digit class (0-9).

    Arguments:
    X (np.ndarray): Feature matrix containing pixel values.
    y (np.ndarray): Labels corresponding to images.
    """

    plt.figure(figsize=(10, 4))
    unique_classes = np.unique(y) # Get unique class labels

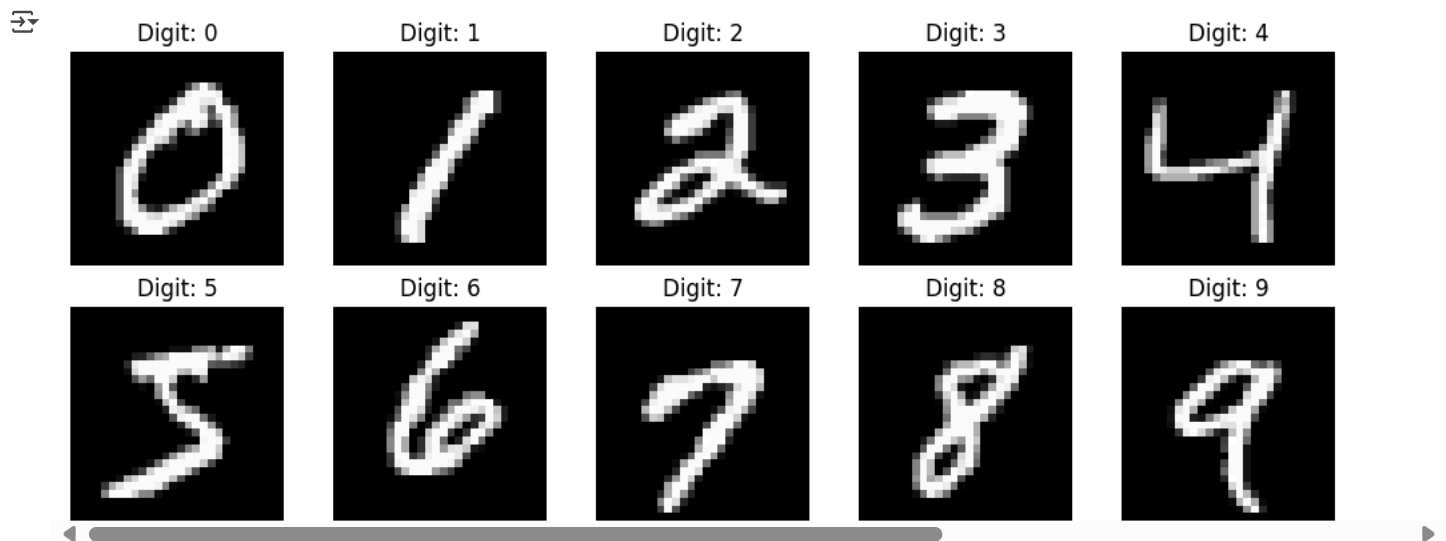
    for i, digit in enumerate(unique_classes):
        index = np.where(y == digit)[0][0] # Find first occurrence of the class
        image = X[index].reshape(28, 28) # Reshape 1D array to 28x28

        plt.subplot(2, 5, i + 1)
        plt.imshow(image, cmap='gray')
        plt.title(f"Digit: {digit}")
        plt.axis('off')

    plt.tight_layout()
    plt.show()

csv_file_path = "/content/drive/MyDrive/Final-Year AI/week2/mnist_dataset.csv" # Path to saved dataset
X_train, X_test, y_train, y_test = load_and_prepare_mnist(csv_file_path)

```



### ✓ A Quick debugging Step:

```

# Assert that X and y have matching lengths
assert len(X_train) == len(y_train), f"Error: X and y have different lengths! X={len(X_train)}, y={len(y_train)}"
print("Move forward: Dimension of Feature Matrix X and label vector y matched.")

```

➡ Move forward: Dimension of Feature Matrix X and label vector y matched.

### ✓ Train the Model:

```
print(f"Training data shape: {X_train.shape}")
print(f"Test data shape: {X_test.shape}")
```

```
↗ Training data shape: (48000, 784)
   Test data shape: (12000, 784)
```

```
from sklearn.preprocessing import OneHotEncoder

# Check if y_train is one-hot encoded
if len(y_train.shape) == 1:
    encoder = OneHotEncoder(sparse_output=False) # Use sparse_output=False for newer versions of sklearn
    y_train = encoder.fit_transform(y_train.reshape(-1, 1)) # One-hot encode labels
    y_test = encoder.transform(y_test.reshape(-1, 1)) # One-hot encode test labels

# Now y_train is one-hot encoded, and we can proceed to use it
d = X_train.shape[1] # Number of features (columns in X_train)
c = y_train.shape[1] # Number of classes (columns in y_train after one-hot encoding)


# Initialize weights with small random values and biases with zeros
W = np.random.randn(d, c) * 0.01 # Small random weights initialized
b = np.zeros(c) # Bias initialized to 0

# Set hyperparameters for gradient descent
alpha = 0.1 # Learning rate
n_iter = 1000 # Number of iterations to run gradient descent

# Train the model using gradient descent
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, y_train, W, b, alpha, n_iter, show_cost=True)

# Plot the cost history to visualize the convergence
plt.plot(cost_history)
plt.title('Cost Function vs. Iterations')
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.grid(True)
plt.show()
```

```
↗ Iteration 0: Cost = 2.197309
   Iteration 100: Cost = 0.607172
   Iteration 200: Cost = 0.489455
   Iteration 300: Cost = 0.440847
   Iteration 400: Cost = 0.412771
   Iteration 500: Cost = 0.393910
   Iteration 600: Cost = 0.380095
   Iteration 700: Cost = 0.369395
   Iteration 800: Cost = 0.360776
   Iteration 900: Cost = 0.353633
```



## ✓ Evaluating the Model:



```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

# Evaluation Function
def evaluate_classification(y_true, y_pred):
    """
    Evaluate classification performance using confusion matrix, precision, recall, and F1-score.

    Parameters:
    y_true (numpy.ndarray): True labels
    y_pred (numpy.ndarray): Predicted labels

    Returns:
    tuple: Confusion matrix, precision, recall, F1 score
    """
    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)

    # Compute precision, recall, and F1-score
    precision = precision_score(y_true, y_pred, average='weighted')
    recall = recall_score(y_true, y_pred, average='weighted')
    f1 = f1_score(y_true, y_pred, average='weighted')

    return cm, precision, recall, f1

# Predict on the test set
y_pred_test = predict_softmax(X_test, W_opt, b_opt)

# Evaluate accuracy
y_test_labels = np.argmax(y_test, axis=1) # True labels in numeric form

# Evaluate the model
cm, precision, recall, f1 = evaluate_classification(y_test_labels, y_pred_test)

# Print the evaluation metrics
print("\nConfusion Matrix:")
print(cm)
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-Score: {f1:.2f}")

# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(12, 12))
cax = ax.imshow(cm, cmap='Blues') # Use a color map for better visualization

# Dynamic number of classes
num_classes = cm.shape[0]
ax.set_xticks(range(num_classes))
ax.set_yticks(range(num_classes))
ax.set_xticklabels([f'Predicted {i}' for i in range(num_classes)])
ax.set_yticklabels([f'Actual {i}' for i in range(num_classes)])

# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if cm[i, j] > np.max(cm) / 2 else 'black')

# Add grid lines and axis labels
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('Actual Label', fontsize=12)

# Adjust layout
plt.tight_layout()
plt.colorbar(cax)
plt.show()

```



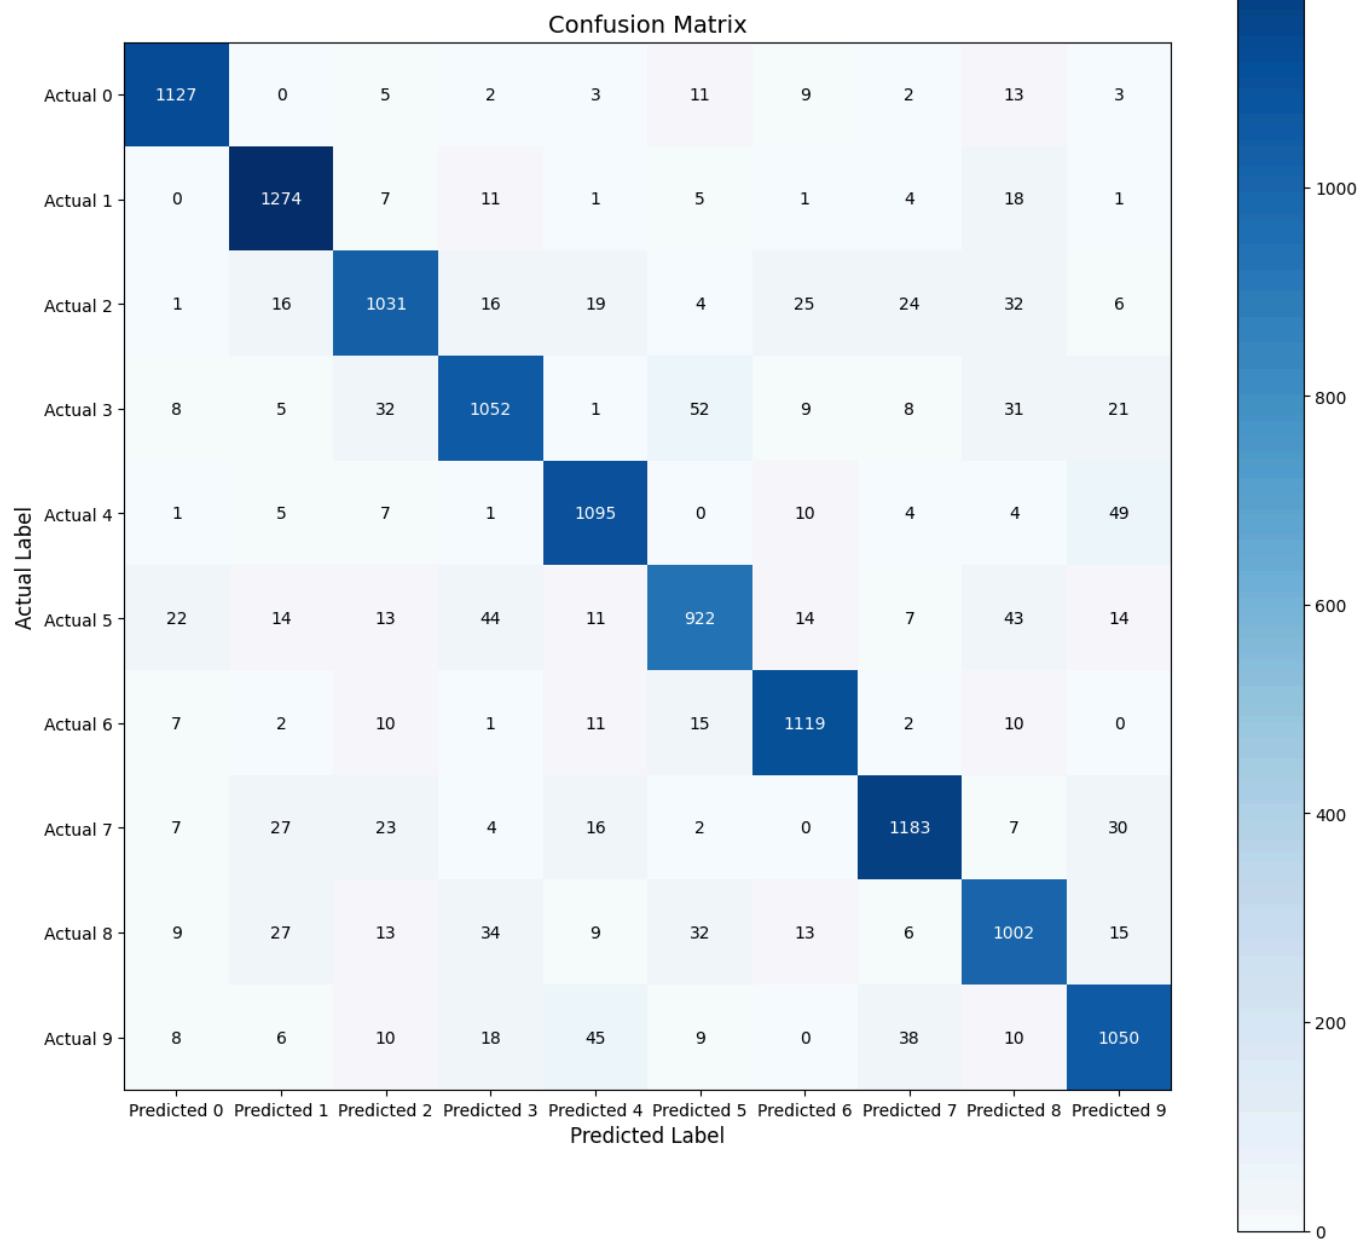
Confusion Matrix:

```
[[1127  0  5  2  3 11  9  2 13  3]
 [  0 1274  7 11  1  5  1  4 18  1]
 [  1 16 1031 16 19  4 25 24 32  6]
 [  8  5 32 1052  1 52  9  8 31 21]
 [  1  5  7  1 1095  0 10  4  4 49]
 [ 22 14 13 44 11 922 14  7 43 14]
 [  7  2 10  1 11 15 1119  2 10  0]
 [  7 27 23  4 16  2  0 1183  7 30]
 [  9 27 13 34  9 32 13  6 1002 15]
 [  8  6 10 18 45  9  0 38 10 1050]]
```

Precision: 0.90

Recall: 0.90

F1-Score: 0.90



## ✓ Linear Separability and Logistic Regression:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification, make_circles
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

# Set random seed for reproducibility
np.random.seed(42)

# Generate a synthetic dataset that is linearly separable
X_linear_separable, y_linear_separable = make_classification(
    n_samples=200, n_features=2, n_informative=2, n_redundant=0, n_clusters_per_class=1, random_state=42
)

# Split the dataset into training and testing sets (80% train, 20% test)
X_train_linear, X_test_linear, y_train_linear, y_test_linear = train_test_split(
    X_linear_separable, y_linear_separable, test_size=0.2, random_state=42
)

# Train a Logistic Regression model on the linearly separable dataset
logistic_model_linear_separable = LogisticRegression()
logistic_model_linear_separable.fit(X_train_linear, y_train_linear)

# Generate a synthetic dataset that is non-linearly separable (circles pattern)
X_non_linear_separable, y_non_linear_separable = make_circles(
    n_samples=200, noise=0.1, factor=0.5, random_state=42
)

# Split the dataset into training and testing sets (80% train, 20% test)
X_train_non_linear, X_test_non_linear, y_train_non_linear, y_test_non_linear = train_test_split(
    X_non_linear_separable, y_non_linear_separable, test_size=0.2, random_state=42
)

# Train a Logistic Regression model on the non-linearly separable dataset
logistic_model_non_linear_separable = LogisticRegression()
logistic_model_non_linear_separable.fit(X_train_non_linear, y_train_non_linear)

# Function to plot the decision boundary of a trained model
def plot_decision_boundary(ax, model, X, y, title):
    h = 0.02 # Step size for mesh grid
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

    # Create a mesh grid over the feature space
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    # Predict class labels for each point in the grid
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot the decision boundary using contour plot
    ax.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)

    # Scatter plot of the actual data points
    ax.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Paired)

    # Formatting the plot
    ax.set_title(title)
    ax.set_xlabel('Feature 1')
    ax.set_ylabel('Feature 2')

# Create subplots to visualize decision boundaries
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Plot decision boundary for linearly separable data (Training set)
plot_decision_boundary(axes[0, 0], logistic_model_linear_separable, X_train_linear, y_train_linear, 'Linearly Separable Data (Training)')

# Plot decision boundary for linearly separable data (Testing set)
plot_decision_boundary(axes[0, 1], logistic_model_linear_separable, X_test_linear, y_test_linear, 'Linearly Separable Data (Testing)')

# Plot decision boundary for non-linearly separable data (Training set)
plot_decision_boundary(axes[1, 0], logistic_model_non_linear_separable, X_train_non_linear, y_train_non_linear, 'Non-Linearly Separable Data (Training)')
```

```
# Plot decision boundary for non-linearly separable data (testing set)
plot_decision_boundary(axes[1, 1], logistic_model_non_linear_separable, X_test_non_linear, y_test_non_linear, 'Non-Linearly Separable Data (T

# Adjust layout for better spacing
plt.tight_layout()

# Save the plot as a PNG file
plt.savefig('testing.png')

# Display the plots
plt.show()
```

